




Chunking gains depend on the model and attack construction

Full input 512-word windows Without redundant tails



False alarms and recorded request time

Model / cohort	Clean flags	Time ratio
E2B / original	0/20 → 1/20	3.04×
E2B / numeric	0/18 → 0/18	1.72×
E4B / original	0/20 → 0/20	2.58×
E4B / numeric	0/18 → 0/18	1.67×

Time: distinct window inputs / full inputs.
Cached-source work counted once; repeated papers aid reuse.

Gemma 4 Q4_K_M • requested high reasoning, 1,024-token cap • preserved 512-word windows, stride 384 • max score >0.5
Correlated, ordered samples; the numeric probe is post-hoc. Counts are descriptive, not deployment accuracy. Time is not energy or uncached latency.